The Geometry of Linear Regression

Lecture Notes: Seeing OLS Through the Projection Theorem

Supplementary notes for the lecture on linear regression

Statistical Machine Learning Lab

Spring 2026

Abstract

These lecture notes develop ordinary least squares (OLS) regression entirely from the perspective of the *Projection Theorem* in inner product spaces. Every classical result—the normal equations, variance decomposition, R^2, the Gauss–Markov theorem, the Frisch–Waugh–Lovell theorem, leverage and influence diagnostics, multicollinearity, omitted variable bias, Ridge and LASSO regularisation, the F-test, and confidence ellipsoids—is derived as a geometric consequence of orthogonal projection onto the column space of the design matrix. The unifying principle: *regression is a shadow cast by the data vector onto the column space, and the residual is the perpendicular distance from object to shadow.*

Contents

1 Prologue: One Theorem, Ten Perspectives **2**

2 Where Coefficients Come From **4**
2.1 Setup and notation . . . 4
2.2 OLS as orthogonal projection . . . 5
2.3 The normal equations . . . 5
2.4 Building intuition: from $\mathbb{R}^2$ to $\mathbb{R}^n$ . . . 5

3 The Hat Matrix and the Right Angle **7**
3.1 The hat matrix . . . 7
3.2 The annihilator matrix . . . 7

4 Pythagoras Runs a Regression **8**
4.1 Variance decomposition as the Pythagorean theorem . . . 8
4.2 R^2 is the cosine of an angle . . . 9
4.3 The Gauss–Markov theorem: OLS is the closest point . . . 11
4.4 R^2 monotonicity . . . 12
4.5 Adjusted R^2: paying for dimensions . . . 13
4.6 Degrees of freedom: the dimension of the residual space . . . 13

5 Peeling Apart the Variables: Frisch–Waugh–Lovell **14**
5.1 What "controlling for" means geometrically . . . 14
5.2 The effect of predictor correlation . . . 18

1 Prologue: One Theorem, Ten Perspectives

New Year's Night, 1802

On the first day of the nineteenth century, the astronomer Giuseppe Piazzi spotted something that should not have been there: a faint, slow-moving point of light in the constellation Taurus. Night after night he tracked it, recording its position against the fixed stars. He had discovered Ceres—the first asteroid, and the largest object between Mars and Jupiter.

Then disaster struck. After only 41 nights of observation, Ceres drifted too close to the Sun and vanished in the glare. Piazzi had captured barely three percent of the asteroid's orbit. As the months passed, the astronomical community faced a crisis: by the time Ceres emerged from behind the Sun, it could be *anywhere* along a vast arc of sky. Several astronomers tried to predict its return, assuming a circular orbit; their predictions would miss by more than ten degrees. The needle had disappeared into a haystack of stars.

A 24-year-old mathematician in Brunswick named Carl Friedrich Gauss took a different approach.

Gauss took Piazzi's 41 data points—each one contaminated by atmospheric distortion, instrument imprecision, and the tremor of a human hand—and asked a question that would reshape the mathematics of uncertainty: *Of all the elliptical orbits that the laws of physics permit, which one passes closest to these noisy observations?*

His method was, at its core, simple. He wrote down the equations of a Keplerian ellipse, parameterised by six orbital elements. These equations carved out a surface—a subspace of possible trajectories—inside the much larger space of all conceivable sequences of celestial positions. Piazzi's 41 observations lived somewhere in that larger space, displaced from the true orbit by the accumulated noise of every measurement. Gauss sought the single point on the orbital surface that lay *nearest* to the data.

He found it by dropping a perpendicular.

The computation was ferocious: weeks of hand calculation, pages of logarithm tables, corrections for the Sun's gravitational tug. But the geometric principle was clear. The best-fitting orbit was the **shadow** of the noisy data cast onto the subspace of physically possible trajectories—the point where a perpendicular from the data meets the surface. The residual—the gap between data and prediction—stood at a right angle to every direction the orbit could vary. Gauss kept his method private; only years later would it become known as the *method of least squares.*

Gauss sent his prediction to the astronomer Franz Xaver von Zach. On December 31, 1801—New Year's Eve—von Zach aimed his telescope at the patch of sky Gauss had specified. There, within half a degree of the predicted position, was Ceres.

It was one of the greatest predictions in the history of science. And it was, at its mathematical heart, an orthogonal projection.

From the stars to the spreadsheet

Two centuries later, the same geometry appears every time you fit a model to data. The setting is humbler—you might be predicting house prices instead of asteroid trajectories—but the underlying act is identical. Your data vector $\boldsymbol{y} \in \mathbb{R}^n$ lives in a high-dimensional space. The column space of your design matrix $\mathcal{C}(\boldsymbol{X})$ is a subspace of possible fitted values, one for each choice of coefficients. The ordinary least squares fit $\hat{\boldsymbol{y}}$ is the point in $\mathcal{C}(\boldsymbol{X})$ closest to $\boldsymbol{y}$. The residual $\boldsymbol{e} = \boldsymbol{y} - \hat{\boldsymbol{y}}$ is perpendicular to the column space.

Gauss projected noisy observations onto the subspace of Keplerian ellipses. You project noisy responses onto the subspace spanned by your regressors. The theorem that makes both possible is the same.

The Flashlight Metaphor

Imagine a flashlight shining straight down onto a table. The **object** is the data vector $\boldsymbol{y} \in \mathbb{R}^n$, hovering above. The **table** is the column space $\mathcal{C}(\boldsymbol{X})$. The **shadow** on the table is the fitted vector $\hat{\boldsymbol{y}}$. The **gap** between object and shadow is the residual $\boldsymbol{e}$. The shadow falls directly below—the residual is perpendicular to the table.

This single image—a shadow cast at right angles—generates everything in these notes.

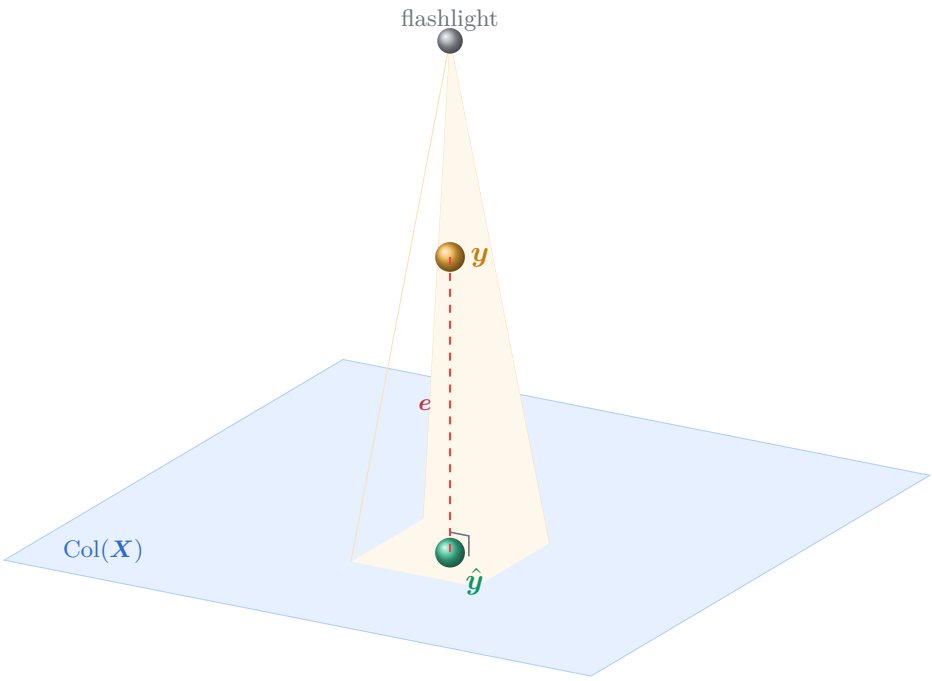

Figure 1: The flashlight metaphor. Think of the data vector $\boldsymbol{y}$ as an object suspended above the column space of $\boldsymbol{X}$. A flashlight shining straight down casts a *shadow* $\hat{\boldsymbol{y}}$ on the plane. The vertical gap between object and shadow is the residual $\boldsymbol{e}$, and it meets the plane at a right angle — exactly as orthogonal projection demands.

Theorem 1.1 (The Projection Theorem—finite-dimensional case)**.** *Let S be a subspace of $\mathbb{R}^n$, and let $\boldsymbol{y} \in \mathbb{R}^n$. There exists a unique $\hat{\boldsymbol{y}} \in S$ such that*

$$\|\boldsymbol{y} - \hat{\boldsymbol{y}}\| \;\leq\; \|\boldsymbol{y} - \boldsymbol{s}\| \qquad \textit{for all } \boldsymbol{s} \in S.$$

Moreover, $\hat{\boldsymbol{y}}$ is characterised by the orthogonality condition

$$(\boldsymbol{y} - \hat{\boldsymbol{y}}) \perp S,$$

i.e. $\langle \boldsymbol{y} - \hat{\boldsymbol{y}}, \boldsymbol{s} \rangle = 0$ for every $\boldsymbol{s} \in S$.

Ordinary least squares is the Projection Theorem applied to $S = \mathcal{C}(\boldsymbol{X})$. That one sentence is the seed from which the rest of these notes grow.

Roadmap

Every section that follows answers a question that the Projection Theorem makes natural:

1. **Where do the coefficients come from?** (§2)
2. **How good is the fit?** (§4)

3. **What does each variable contribute on its own?** (§5)
4. **What can go wrong, and how do we fix it?** (§??–§??)
5. **How do we test whether the shadow is real?** (§??)

By the time you reach the end, you will see that Gauss's prediction of Ceres, the R^2 on your next problem set, and the p-value in a clinical trial are all manifestations of the same geometric act: dropping a perpendicular onto a subspace, and measuring the gap that remains.

The Projection Theorem guarantees that a closest point exists and is characterised by orthogonality. But how do we actually compute *it? The next section translates the geometric existence result into an explicit formula—the normal equations—and shows that the OLS coefficients are simply the coordinates of this closest point.*

? Stop and Think

Exercise. Take $\boldsymbol{y} = (3, 1)^\top$ and $\boldsymbol{x} = (2, 1)^\top$ in $\mathbb{R}^2$. The column space $\mathcal{C}(\boldsymbol{x})$ is the line through the origin in the direction $(2, 1)$.

1. Sketch the vectors $\boldsymbol{y}$ and $\boldsymbol{x}$ in $\mathbb{R}^2$.
2. Compute the projection of $\boldsymbol{y}$ onto $\mathcal{C}(\boldsymbol{x})$ using the formula $\hat{\boldsymbol{y}} = \boldsymbol{x}(\boldsymbol{x}^\top\boldsymbol{x})^{-1}\boldsymbol{x}^\top\boldsymbol{y}$. Where does the projection land on the line?
3. Draw the residual vector $\boldsymbol{e} = \boldsymbol{y} - \hat{\boldsymbol{y}}$. Verify visually that it is perpendicular to $\boldsymbol{x}$.
4. Compute $\boldsymbol{x}^\top\boldsymbol{e}$ and confirm it equals zero.

Answer check: You should get $\hat{\boldsymbol{y}} = \frac{7}{5}(2, 1)^\top = (2.8,\ 1.4)^\top$ and $\boldsymbol{e} = (0.2,\ -0.4)^\top$.

2 Where Coefficients Come From

Standing Assumption

Full column rank. Throughout these notes (unless stated otherwise), we assume that the $n \times p$ design matrix $\boldsymbol{X}$ has **full column rank**: $\operatorname{rank}(\boldsymbol{X}) = p$, with $n \geq p$.
This is equivalent to any of the following:

(i) The columns of $\boldsymbol{X}$ are linearly independent.
(ii) The Gram matrix $\boldsymbol{X}^\top\boldsymbol{X} \in \mathbb{R}^{p\times p}$ is invertible (positive definite).
(iii) The column space $\mathcal{C}(\boldsymbol{X})$ has dimension exactly p.
(iv) No predictor is an exact linear combination of the others.

Without this assumption, $(\boldsymbol{X}^\top\boldsymbol{X})^{-1}$ does not exist, and the OLS coefficient vector $\hat{\boldsymbol{\beta}}$ is not uniquely determined (although the fitted values $\hat{\boldsymbol{y}} = \boldsymbol{H}\boldsymbol{y}$ remain unique as a projection).

2.1 Setup and notation

We observe n data points (x_i, y_i), $i = 1, \ldots, n$. Stack the responses into a single vector $\boldsymbol{y} = (y_1, \ldots, y_n)^\top \in \mathbb{R}^n$ and form the $n \times p$ **design matrix**

$$\boldsymbol{X} = \begin{pmatrix} \boldsymbol{x}_1^\top \\ \vdots \\ \boldsymbol{x}_n^\top \end{pmatrix} \in \mathbb{R}^{n\times p},$$

where typically the first column is a column of ones (the intercept). The **column space** of $\boldsymbol{X}$ is

$$\mathcal{C}(\boldsymbol{X}) = \{\boldsymbol{X}\boldsymbol{\beta} : \boldsymbol{\beta} \in \mathbb{R}^p\} \subseteq \mathbb{R}^n.$$

Every vector in $\mathcal{C}(\boldsymbol{X})$ represents the set of fitted values for *some* choice of coefficients.

2.2 OLS as orthogonal projection

Definition 2.1 (OLS estimator)**.** The OLS estimator is

$$\hat{\boldsymbol{\beta}} = \arg\min_{\boldsymbol{\beta}\in\mathbb{R}^p} \|\boldsymbol{y} - \boldsymbol{X}\boldsymbol{\beta}\|^2 .$$

Since $\mathcal{C}(\boldsymbol{X})$ is a closed subspace of $\mathbb{R}^n$, the Projection Theorem (1.1) guarantees a unique minimiser. The fitted values $\hat{\boldsymbol{y}} = \boldsymbol{X}\hat{\boldsymbol{\beta}}$ are the orthogonal projection of $\boldsymbol{y}$ onto $\mathcal{C}(\boldsymbol{X})$, and the residual $\boldsymbol{e} = \boldsymbol{y} - \hat{\boldsymbol{y}}$ is perpendicular to $\mathcal{C}(\boldsymbol{X})$.

Key Geometric Insight

$\hat{\boldsymbol{\beta}}$ is the coordinate vector of the foot of the perpendicular from $\boldsymbol{y}$ to $\mathcal{C}(\boldsymbol{X})$. The coefficients are not arbitrary numbers from a formula—they locate the point on the column space closest to the data.

2.3 The normal equations

The orthogonality condition $\boldsymbol{e} \perp \mathcal{C}(\boldsymbol{X})$ means $\langle \boldsymbol{X}\boldsymbol{v}, \boldsymbol{e}\rangle = 0$ for all $\boldsymbol{v} \in \mathbb{R}^p$, i.e.

$$\boldsymbol{X}^\top \boldsymbol{e} = \boldsymbol{0} \qquad \Longleftrightarrow \qquad \boldsymbol{X}^\top(\boldsymbol{y} - \boldsymbol{X}\hat{\boldsymbol{\beta}}) = \boldsymbol{0} \qquad \Longleftrightarrow \qquad \boxed{\boldsymbol{X}^\top \boldsymbol{X}\,\hat{\boldsymbol{\beta}} = \boldsymbol{X}^\top \boldsymbol{y}.}$$

These are the **normal equations**. If $\boldsymbol{X}$ has full column rank, $\boldsymbol{X}^\top\boldsymbol{X}$ is invertible and

$$\hat{\boldsymbol{\beta}} = (\boldsymbol{X}^\top\boldsymbol{X})^{-1}\boldsymbol{X}^\top\boldsymbol{y}. \tag{1}$$

The formula is a recipe for dropping a perpendicular.

2.4 Building intuition: from $\mathbb{R}^2$ to $\mathbb{R}^n$

- With $n = 2$ observations, $\boldsymbol{y} \in \mathbb{R}^2$ and the column space is a line through the origin (for simple regression with intercept, a line in the plane).
- With $n = 3$, $\boldsymbol{y} \in \mathbb{R}^3$ and a model with intercept and one predictor spans a 2-dimensional plane—the geometry is fully visible.
- For $n = 2000$, the geometry is identical; only the ambient dimension changes. The right angle between $\boldsymbol{e}$ and $\mathcal{C}(\boldsymbol{X})$ holds in every dimension.

Example 2.2 (Your First Regression in $\mathbb{R}^3$)**.** We have $n = 3$ observations with one predictor x and a response y:

i	x_i	y_i
1	1	1
2	2	3
3	3	2

The model is $y_i = \beta_0 + \beta_1 x_i + e_i$, so the design matrix and response vector are

$$\boldsymbol{X} = \begin{pmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \end{pmatrix}, \qquad \boldsymbol{y} = \begin{pmatrix} 1 \\ 3 \\ 2 \end{pmatrix}.$$

Since $\boldsymbol{y} \in \mathbb{R}^3$ and $\mathcal{C}(\boldsymbol{X})$ is a 2-dimensional plane in $\mathbb{R}^3$, this is literally a point–plane–shadow problem.

Step 1: Compute $\boldsymbol{X}^\top\boldsymbol{X}$ and $\boldsymbol{X}^\top\boldsymbol{y}$.

$$\boldsymbol{X}^\top\boldsymbol{X} = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix}\begin{pmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \end{pmatrix} = \begin{pmatrix} 3 & 6 \\ 6 & 14 \end{pmatrix}, \qquad \boldsymbol{X}^\top\boldsymbol{y} = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix}\begin{pmatrix} 1 \\ 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 6 \\ 13 \end{pmatrix}.$$

Step 2: Solve $\boldsymbol{X}^\top \boldsymbol{X}\,\hat{\boldsymbol{\beta}} = \boldsymbol{X}^\top \boldsymbol{y}$.

The determinant is $\det(\boldsymbol{X}^\top \boldsymbol{X}) = 3 \cdot 14 - 6 \cdot 6 = 42 - 36 = 6$, so

$$(\boldsymbol{X}^\top \boldsymbol{X})^{-1} = \frac{1}{6}\begin{pmatrix} 14 & -6 \\ -6 & 3 \end{pmatrix} = \begin{pmatrix} 7/3 & -1 \\ -1 & 1/2 \end{pmatrix}.$$

Therefore

$$\hat{\boldsymbol{\beta}} = (\boldsymbol{X}^\top \boldsymbol{X})^{-1}\boldsymbol{X}^\top \boldsymbol{y} = \begin{pmatrix} 7/3 & -1 \\ -1 & 1/2 \end{pmatrix}\begin{pmatrix} 6 \\ 13 \end{pmatrix} = \begin{pmatrix} 14 - 13 \\ -6 + 13/2 \end{pmatrix} = \begin{pmatrix} 1 \\ 1/2 \end{pmatrix}.$$

The fitted line is $\hat{y} = 1 + \frac{1}{2}x$.

Step 3: Compute $\hat{\boldsymbol{y}}$ and $\boldsymbol{e}$.

$$\hat{\boldsymbol{y}} = \boldsymbol{X}\hat{\boldsymbol{\beta}} = \begin{pmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \end{pmatrix}\begin{pmatrix} 1 \\ 1/2 \end{pmatrix} = \begin{pmatrix} 3/2 \\ 2 \\ 5/2 \end{pmatrix}, \qquad \boldsymbol{e} = \boldsymbol{y} - \hat{\boldsymbol{y}} = \begin{pmatrix} 1 \\ 3 \\ 2 \end{pmatrix} - \begin{pmatrix} 3/2 \\ 2 \\ 5/2 \end{pmatrix} = \begin{pmatrix} -1/2 \\ 1 \\ -1/2 \end{pmatrix}.$$

Step 4: Verify the right angle ($\boldsymbol{X}^\top \boldsymbol{e} = \boldsymbol{0}$).

$$\boldsymbol{X}^\top \boldsymbol{e} = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 2 & 3 \end{pmatrix}\begin{pmatrix} -1/2 \\ 1 \\ -1/2 \end{pmatrix} = \begin{pmatrix} -1/2 + 1 - 1/2 \\ -1/2 + 2 - 3/2 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}. \checkmark$$

The residual vector is perpendicular to *every* column of $\boldsymbol{X}$, hence perpendicular to the entire column space—exactly as the Projection Theorem demands.

The geometry. The data vector $\boldsymbol{y} = (1, 3, 2)^\top$ is a point in $\mathbb{R}^3$. The column space $\mathcal{C}(\boldsymbol{X})$ is the plane spanned by $(1, 1, 1)^\top$ and $(1, 2, 3)^\top$. The fitted vector $\hat{\boldsymbol{y}} = (3/2,\ 2,\ 5/2)^\top$ is the foot of the perpendicular from $\boldsymbol{y}$ to this plane. The residual $\boldsymbol{e} = (-1/2,\ 1,\ -1/2)^\top$ is the perpendicular itself. We verified that $\boldsymbol{e}$ is orthogonal to the plane.

Takeaway. OLS literally drops a perpendicular from $\boldsymbol{y}$ to $\mathcal{C}(\boldsymbol{X})$; the normal equations $\boldsymbol{X}^\top \boldsymbol{e} = \boldsymbol{0}$ are the *right-angle condition*, and the hat matrix is the machine that does it.

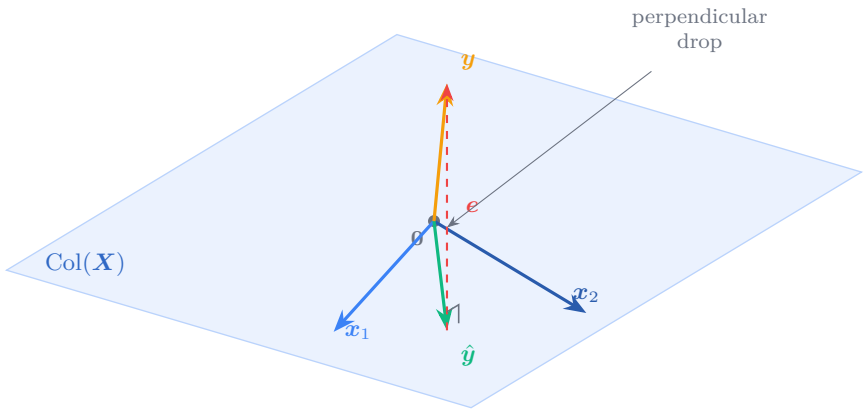

Figure 2: Projection in $\mathbb{R}^3$. The column space $\mathrm{Col}(\boldsymbol{X})$ is the plane spanned by $\boldsymbol{x}_1$ and $\boldsymbol{x}_2$. The fitted vector $\hat{\boldsymbol{y}}$ is the point in this plane closest to $\boldsymbol{y}$, obtained by dropping a perpendicular. The residual $\boldsymbol{e} = \boldsymbol{y} - \hat{\boldsymbol{y}}$ is orthogonal to every vector in the plane, which is precisely the content of the normal equations $\boldsymbol{X}^\top \boldsymbol{e} = \boldsymbol{0}$.

We now have a formula for the projection and its coordinates. But the formula hides a remarkable structure: the projection is a linear operator $\boldsymbol{H}$ *with powerful algebraic properties. What can these properties tell us about individual observations, and about the projection itself?*

3 The Hat Matrix and the Right Angle

3.1 The hat matrix

Definition 3.1. The **hat matrix** (or projection matrix) is

$$\boldsymbol{H} = \boldsymbol{X}(\boldsymbol{X}^\top \boldsymbol{X})^{-1}\boldsymbol{X}^\top.$$

It "puts the hat on $\boldsymbol{y}$": $\hat{\boldsymbol{y}} = \boldsymbol{H}\boldsymbol{y}$.

Proposition 3.2 (Properties of $\boldsymbol{H}$)**.**
(i) ***Idempotent:*** $\boldsymbol{H}^2 = \boldsymbol{H}$.
(ii) ***Symmetric:*** $\boldsymbol{H}^\top = \boldsymbol{H}$.
(iii) $\operatorname{rank}(\boldsymbol{H}) = \operatorname{tr}(\boldsymbol{H}) = p$.
(iv) $0 \le h_{ii} \le 1$ *for every diagonal entry, and* $\sum_i h_{ii} = p$.

Proof. (i) $\boldsymbol{H}^2 = \boldsymbol{X}(\boldsymbol{X}^\top\boldsymbol{X})^{-1}\boldsymbol{X}^\top\boldsymbol{X}(\boldsymbol{X}^\top\boldsymbol{X})^{-1}\boldsymbol{X}^\top = \boldsymbol{X}(\boldsymbol{X}^\top\boldsymbol{X})^{-1}\boldsymbol{X}^\top = \boldsymbol{H}$. (ii) Follows directly from the symmetry of $(\boldsymbol{X}^\top\boldsymbol{X})^{-1}$. (iii)–(iv) follow from idempotency: the eigenvalues of an idempotent matrix are 0 or 1, so $\operatorname{tr}(\boldsymbol{H}) = \operatorname{rank}(\boldsymbol{H}) = p$, and each diagonal entry $h_{ii} = \boldsymbol{e}_i^\top \boldsymbol{H}\boldsymbol{e}_i$ lies in $[0, 1]$ because $\boldsymbol{H}$ is positive semi-definite with eigenvalues in $\{0, 1\}$. □

Geometric meaning. Idempotency says: *the shadow of a shadow is itself.* If $\hat{\boldsymbol{y}}$ is already on the column space, projecting it again changes nothing. Symmetry says the projection is *orthogonal* (perpendicular), not oblique. A non-symmetric idempotent matrix would still decompose $\boldsymbol{y} = \boldsymbol{P}\boldsymbol{y} + (\boldsymbol{I} - \boldsymbol{P})\boldsymbol{y}$ into a component in $\mathcal{C}(\boldsymbol{P})$ and a remainder, but the two components would not be perpendicular—the "shadow" would fall at a slant. Symmetry of $\boldsymbol{H}$ is what guarantees that the residual $\boldsymbol{e} = (\boldsymbol{I} - \boldsymbol{H})\boldsymbol{y}$ is truly perpendicular to the column space, so that $\langle \hat{\boldsymbol{y}}, \boldsymbol{e}\rangle = \boldsymbol{y}^\top\boldsymbol{H}^\top(\boldsymbol{I} - \boldsymbol{H})\boldsymbol{y} = \boldsymbol{y}^\top(\boldsymbol{H} - \boldsymbol{H}^2)\boldsymbol{y} = 0$. In short: *idempotency makes it a projection; symmetry makes it orthogonal.*

Remark 3.3 (Characterisation of orthogonal projections)**.** A fundamental theorem (proved in Appendix **??**) states that a matrix is an orthogonal projection onto its column space if and only if it is symmetric and idempotent.

- The hat matrix $\boldsymbol{H} = \boldsymbol{X}(\boldsymbol{X}^\top\boldsymbol{X})^{-1}\boldsymbol{X}^\top$ satisfies both properties (Proposition 3.2), confirming it is the orthogonal projection onto $\mathcal{C}(\boldsymbol{X})$.
- The annihilator $\boldsymbol{M} = \boldsymbol{I} - \boldsymbol{H}$ likewise satisfies both, confirming it is the orthogonal projection onto $\mathcal{C}(\boldsymbol{X})^\perp$.

3.2 The annihilator matrix

Define $\boldsymbol{M} = \boldsymbol{I}_n - \boldsymbol{H}$. Then $\boldsymbol{e} = \boldsymbol{M}\boldsymbol{y}$, and $\boldsymbol{M}$ inherits the same properties: $\boldsymbol{M}^2 = \boldsymbol{M}$, $\boldsymbol{M}^\top = \boldsymbol{M}$. It is the projection onto the *orthogonal complement* $\mathcal{C}(\boldsymbol{X})^\perp$. Together, $\boldsymbol{H}$ and $\boldsymbol{M}$ decompose any vector into two perpendicular components:

$$\boldsymbol{y} = \underbrace{\boldsymbol{H}\boldsymbol{y}}_{\hat{\boldsymbol{y}} \in \mathcal{C}(\boldsymbol{X})} + \underbrace{\boldsymbol{M}\boldsymbol{y}}_{\boldsymbol{e} \in \mathcal{C}(\boldsymbol{X})^\perp}.$$

Warning

The orthogonality $\boldsymbol{X}^\top\boldsymbol{e} = \boldsymbol{0}$ is a property of the *method*, not a validation of the *model*. You could regress salary on zodiac sign and the residuals would still be perpendicular to the zodiac column. Residual-predictor uncorrelatedness is by construction.

Misconception Alert

"The projection validates the model."
A common student error: "My residuals are uncorrelated with my predictors, so the model must be correct." This confuses a *property of the method* with *evidence for the model.* The orthogonality $\boldsymbol{X}^\top \boldsymbol{e} = \boldsymbol{0}$ holds for *any* design matrix $\boldsymbol{X}$, regardless of whether $\boldsymbol{X}$ contains the right variables, whether the relationship is linear, or whether the errors are well-behaved. You could regress stock returns on the phases of the moon and the residuals would still be perpendicular to the moon-phase column.
Residual-predictor orthogonality tells you that OLS did its job (found the closest point). It says nothing about whether the column space was the right target.

The hat matrix decomposes $\boldsymbol{y}$ into two perpendicular pieces: the shadow $\hat{\boldsymbol{y}}$ and the residual $\boldsymbol{e}$. Since these pieces are perpendicular, the Pythagorean theorem must apply. What does Pythagoras tell us about regression? The answer is the variance decomposition—and R^2 as an angle.

4 Pythagoras Runs a Regression

4.1 Variance decomposition as the Pythagorean theorem

Since $\hat{\boldsymbol{y}} \perp \boldsymbol{e}$ and $\boldsymbol{y} = \hat{\boldsymbol{y}} + \boldsymbol{e}$, the Pythagorean theorem gives

$$\|\boldsymbol{y}\|^2 = \|\hat{\boldsymbol{y}}\|^2 + \|\boldsymbol{e}\|^2 . \tag{2}$$

Working with demeaned vectors $\tilde{\boldsymbol{y}} = \boldsymbol{y} - \bar{y}\boldsymbol{1}$ (which is the projection of $\boldsymbol{y}$ onto the orthogonal complement of the intercept), we obtain the fundamental variance decomposition. Before stating it, let us define the three sums of squares.

Definition 4.1 (Sums of Squares)**.**

- **SST** (Total Sum of Squares) $= \sum_{i=1}^n (y_i - \bar{y})^2 = \|\tilde{\boldsymbol{y}}\|^2$.
  This measures the *total variability* in the response. It answers: "How spread out are the y_i values around their mean?"
- **SSR** (Regression Sum of Squares, also called the "explained" sum of squares) $= \sum_{i=1}^n (\hat{y}_i - \bar{y})^2 = \left\|\tilde{\hat{\boldsymbol{y}}}\right\|^2$.
  This measures the variability *explained by the model.* It answers: "How much of the spread in $\boldsymbol{y}$ is captured by the fitted values?"
- **SSE** (Error Sum of Squares, also called the "residual" sum of squares) $= \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \|\boldsymbol{e}\|^2$.
  This measures the variability *left unexplained* by the model. It answers: "How far are the actual values from the predictions?"

With these definitions, the Pythagorean theorem gives us:

$$\boxed{\underbrace{\|\tilde{\boldsymbol{y}}\|^2}_{\text{SST}} = \underbrace{\left\|\tilde{\hat{\boldsymbol{y}}}\right\|^2}_{\text{SSR}} + \underbrace{\|\boldsymbol{e}\|^2}_{\text{SSE}}.}$$

In words: the total variability in the data equals the variability explained by the regression plus the variability left over in the residuals. The variance decomposition *is* the Pythagorean theorem applied to the right triangle $(\tilde{\boldsymbol{y}},\ \tilde{\hat{\boldsymbol{y}}},\ \boldsymbol{e})$ in $\mathbb{R}^n$.

Remark 4.2 (The intercept is needed for SST = SSR + SSE)**.** The decomposition SST = SSR + SSE requires that the model includes an intercept (equivalently, that $\boldsymbol{1} \in \mathcal{C}(\boldsymbol{X})$). Here is why.

The decomposition uses demeaned vectors: $\tilde{\boldsymbol{y}} = \boldsymbol{y} - \bar{y}\mathbf{1}$ and $\tilde{\hat{\boldsymbol{y}}} = \hat{\boldsymbol{y}} - \bar{y}\mathbf{1}$. The Pythagorean theorem gives $\|\tilde{\boldsymbol{y}}\|^2 = \left\|\tilde{\hat{\boldsymbol{y}}}\right\|^2 + \|\boldsymbol{e}\|^2$ only if $\tilde{\hat{\boldsymbol{y}}} \perp \boldsymbol{e}$.

Now $\tilde{\hat{\boldsymbol{y}}} = \hat{\boldsymbol{y}} - \bar{y}\mathbf{1}$. We know $\hat{\boldsymbol{y}} \perp \boldsymbol{e}$ is *not* generally true (it holds only when $\boldsymbol{y} \in \mathcal{C}(\boldsymbol{X})$). What we need is $(\hat{\boldsymbol{y}} - \bar{y}\mathbf{1}) \perp \boldsymbol{e}$, which expands to $\hat{\boldsymbol{y}}^\top \boldsymbol{e} = \bar{y}\mathbf{1}^\top \boldsymbol{e}$.

We always have $\hat{\boldsymbol{y}}^\top \boldsymbol{e} = 0$ (since $\hat{\boldsymbol{y}} \in \mathcal{C}(\boldsymbol{X})$ and $\boldsymbol{e} \perp \mathcal{C}(\boldsymbol{X})$). So we need $\bar{y}\mathbf{1}^\top \boldsymbol{e} = 0$. If $\bar{y} \neq 0$, this requires $\mathbf{1}^\top \boldsymbol{e} = 0$, i.e., that the residuals sum to zero.

The residuals sum to zero if and only if $\mathbf{1} \in \mathcal{C}(\boldsymbol{X})$ (because $\boldsymbol{e} \perp \mathcal{C}(\boldsymbol{X})$ and $\mathbf{1} \in \mathcal{C}(\boldsymbol{X})$ implies $\mathbf{1}^\top \boldsymbol{e} = 0$). Without an intercept, residuals need not sum to zero, and the clean decomposition $\mathrm{SST} = \mathrm{SSR} + \mathrm{SSE}$ can fail.

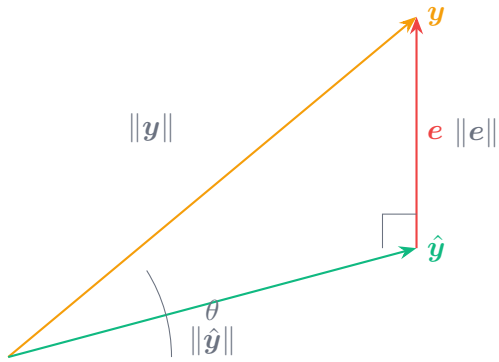

Figure 3: The regression right triangle. The Pythagorean theorem gives $\mathrm{SST} = \mathrm{SSR} + \mathrm{SSE}$.

4.2 R^2 is the cosine of an angle

The coefficient of determination R^2 is usually introduced via the formula $R^2 = 1 - \mathrm{SSE}/\mathrm{SST}$, or equivalently $R^2 = \mathrm{SSR}/\mathrm{SST}$. Let us derive this from the geometry of the regression right triangle.

Look at the right triangle in the figure above. The *hypotenuse* is the demeaned data vector $\tilde{\boldsymbol{y}}$ (with length $\sqrt{\mathrm{SST}}$), the *adjacent side* is the demeaned fitted vector $\tilde{\hat{\boldsymbol{y}}}$ (with length $\sqrt{\mathrm{SSR}}$), and the *opposite side* is the residual $\boldsymbol{e}$ (with length $\sqrt{\mathrm{SSE}}$). The angle θ sits at the origin, between the hypotenuse and the adjacent side.

From basic trigonometry, the cosine of an angle in a right triangle is

$$\cos\theta = \frac{\text{adjacent}}{\text{hypotenuse}} = \frac{\left\|\tilde{\hat{\boldsymbol{y}}}\right\|}{\|\tilde{\boldsymbol{y}}\|} = \frac{\sqrt{\mathrm{SSR}}}{\sqrt{\mathrm{SST}}}.$$

Squaring both sides:

$$\cos^2\theta = \frac{\left\|\tilde{\hat{\boldsymbol{y}}}\right\|^2}{\|\tilde{\boldsymbol{y}}\|^2} = \frac{\mathrm{SSR}}{\mathrm{SST}}.$$

But SSR/SST is precisely the definition of R^2, so we arrive at

$$\boxed{R^2 = \cos^2\theta = \frac{\mathrm{SSR}}{\mathrm{SST}} = 1 - \frac{\mathrm{SSE}}{\mathrm{SST}}.} \tag{3}$$

The last equality follows directly from $\mathrm{SST} = \mathrm{SSR} + \mathrm{SSE}$.

Key Geometric Insight

R^2 is literally the squared cosine of the angle between the data vector and its projection.

- When $\theta \approx 0$, the data vector nearly lies on the column space. The shadow almost coincides with the object, meaning the model explains nearly all the variability: $R^2 \approx 1$.
- When $\theta \approx 90$ř, the data vector is nearly perpendicular to the column space. The shadow is tiny—the model explains almost nothing: $R^2 \approx 0$.
- $R^2 = 0.75$ means $\cos^2\theta = 0.75$, so $\theta \approx 30$ř. Three-quarters of the data's variability is captured by the projection; one quarter is residual.

Example 4.3 (Pythagoras Checks Out)**.** We continue with the data from Example 2.2: $\boldsymbol{y} = (1, 3, 2)^\top$, $\hat{\boldsymbol{y}} = (3/2,\ 2,\ 5/2)^\top$, $\boldsymbol{e} = (-1/2,\ 1,\ -1/2)^\top$.

The sample mean is $\bar{y} = (1 + 3 + 2)/3 = 2$, so the mean vector is $\bar{y}\mathbf{1} = (2, 2, 2)^\top$.

Step 1: Compute the three sums of squares.

Total sum of squares (deviation of $\boldsymbol{y}$ from its mean):

$$\mathrm{SST} = \|\boldsymbol{y} - \bar{y}\mathbf{1}\|^2 = (1-2)^2 + (3-2)^2 + (2-2)^2 = 1 + 1 + 0 = 2.$$

Regression sum of squares (deviation of $\hat{\boldsymbol{y}}$ from the mean):

$$\mathrm{SSR} = \|\hat{\boldsymbol{y}} - \bar{y}\mathbf{1}\|^2 = (3/2-2)^2 + (2-2)^2 + (5/2-2)^2 = 1/4 + 0 + 1/4 = 1/2.$$

Residual sum of squares (length of the residual):

$$\mathrm{SSE} = \|\boldsymbol{e}\|^2 = (-1/2)^2 + 1^2 + (-1/2)^2 = 1/4 + 1 + 1/4 = 3/2.$$

Step 2: Verify Pythagoras.

$$\mathrm{SSR} + \mathrm{SSE} = \frac{1}{2} + \frac{3}{2} = 2 = \mathrm{SST}. \checkmark$$

This is the Pythagorean theorem in $\mathbb{R}^3$: the squared hypotenuse $\|\boldsymbol{y} - \bar{y}\mathbf{1}\|^2$ equals the sum of the two squared legs $\|\hat{\boldsymbol{y}} - \bar{y}\mathbf{1}\|^2 + \|\boldsymbol{e}\|^2$. This works because $\hat{\boldsymbol{y}} - \bar{y}\mathbf{1}$ and $\boldsymbol{e}$ are orthogonal—they live in perpendicular subspaces.

Step 3: Compute R^2 as $\cos^2\theta$.

$$R^2 = \frac{\mathrm{SSR}}{\mathrm{SST}} = \frac{1/2}{2} = \frac{1}{4} = 0.25.$$

Since $R^2 = \cos^2\theta$ where θ is the angle between $\boldsymbol{y} - \bar{y}\mathbf{1}$ and $\hat{\boldsymbol{y}} - \bar{y}\mathbf{1}$:

$$\cos\theta = \sqrt{R^2} = \sqrt{1/4} = 1/2, \qquad \theta = \arccos(1/2) = 60\text{ř}.$$

Interpretation. The angle between the centered data vector and its shadow is 60ř—rather far from 0ř (perfect fit) and closer to 90ř (no fit at all). The model explains only 25% of the variance. Geometrically, the shadow on the regression plane captures only a quarter of the squared length of the centered data vector; the remaining three-quarters is in the perpendicular residual direction.

Takeaway. R^2 is literally the cosine-squared of the angle between the centered observation vector and its projection. A small angle means a long shadow and a good fit; a large angle means the data point hovers far above the column space.

The shadow-sharpness analogy. Think of R^2 as measuring how *sharp* a shadow is. When the flashlight is nearly aligned with the data vector ($\theta \approx 0$), the shadow $\hat{\boldsymbol{y}}$ is almost as long as the object $\boldsymbol{y}$—the shadow is a crisp, faithful reproduction of the object ($R^2 \approx 1$).

When the flashlight is nearly perpendicular to the data ($\theta \approx 90°$), the shadow shrinks to almost nothing—a tiny smudge that preserves none of the object's structure ($R^2 \approx 0$).

A shadow with $R^2 = 0.5$ (i.e. $\theta = 45°$) preserves half the object's "energy" (squared length). The shadow exists, but it is a washed-out version of the original. Whether that is good enough depends on what you need the shadow for.

Misconception Alert

"$R^2 = 0.8$ means the model is good."
R^2 measures a geometric angle—nothing more. A high R^2 does *not* mean:

- The model is correctly specified (you might have the wrong functional form or omitted variables).
- The model will predict well out of sample (overfitting inflates R^2 on training data).
- The coefficients are meaningful (multicollinearity can give $R^2 = 0.99$ with wildly unstable individual coefficients).

Conversely, a low R^2 does *not* mean the model is useless. In fields with inherently noisy data (e.g. daily stock returns), $R^2 = 0.03$ can represent an economically significant and statistically reliable effect.
R^2 answers exactly one question: "What fraction of the total variability in $\boldsymbol{y}$ is captured by the projection onto $\mathcal{C}(\boldsymbol{X})$?" Whether that fraction is "good enough" depends on the scientific context, not on any universal threshold.

? Stop and Think

Thought experiment: adding a random noise column.
Suppose your current model has $R^2 = 0.60$. You add a new column $\boldsymbol{z}$ to $\boldsymbol{X}$, where $\boldsymbol{z}$ is pure random noise (independent of everything).

1. Will SSE increase, decrease, or stay the same? *Hint:* adding a column enlarges $\mathcal{C}(\boldsymbol{X})$.
2. Will SSR increase, decrease, or stay the same?
3. Will R^2 increase, decrease, or stay the same?
4. If you keep adding random noise columns until $p = n$, what happens to R^2?

Answers: (1) SSE can only decrease (or stay the same)—the projection onto a larger subspace is always at least as close. (2) Since SST is fixed and SSE decreases, SSR must increase. (3) R^2 increases—even though the new column is pure noise! (4) When $p = n$, $\mathcal{C}(\boldsymbol{X}) = \mathbb{R}^n$, so $\hat{\boldsymbol{y}} = \boldsymbol{y}$, $\boldsymbol{e} = \boldsymbol{0}$, and $R^2 = 1$. The model "explains" everything but predicts nothing. This is why R^2 alone is not a reliable measure of model quality.

4.3 The Gauss–Markov theorem: OLS is the closest point

Theorem 4.4 (Gauss–Markov)**.** *Assume $E[\boldsymbol{\varepsilon}] = \boldsymbol{0}$ and $Var(\boldsymbol{\varepsilon}) = \sigma^2 \boldsymbol{I}_n$ (no assumption on the distribution of $\boldsymbol{\varepsilon}$). Among all linear unbiased estimators of $\boldsymbol{c}^\top \boldsymbol{\beta}$ (for any fixed $\boldsymbol{c} \in \mathbb{R}^p$), the OLS estimator $\boldsymbol{c}^\top \hat{\boldsymbol{\beta}} = \boldsymbol{c}^\top (\boldsymbol{X}^\top \boldsymbol{X})^{-1} \boldsymbol{X}^\top \boldsymbol{y}$ has the smallest variance. In the vector form: for any linear unbiased estimator $\tilde{\boldsymbol{\beta}} = \boldsymbol{C}\boldsymbol{y}$ of $\boldsymbol{\beta}$,*

$$Var(\tilde{\boldsymbol{\beta}}) - Var(\hat{\boldsymbol{\beta}}) \succeq \boldsymbol{0} \qquad \textit{(positive semi-definite).}$$

Proof. **Step 1: Set up the alternative estimator.** Let $\tilde{\boldsymbol{\beta}} = \boldsymbol{C}\boldsymbol{y}$ be any linear estimator of $\boldsymbol{\beta}$, where $\boldsymbol{C} \in \mathbb{R}^{p \times n}$. Write

$$\boldsymbol{C} = (\boldsymbol{X}^\top \boldsymbol{X})^{-1} \boldsymbol{X}^\top + \boldsymbol{D},$$

where $\boldsymbol{D} = \boldsymbol{C} - (\boldsymbol{X}^\top \boldsymbol{X})^{-1}\boldsymbol{X}^\top$ is the "excess" beyond the OLS matrix.

Step 2: Derive the constraint from unbiasedness. For $\tilde{\boldsymbol{\beta}}$ to be unbiased for $\boldsymbol{\beta}$, we need $E[\tilde{\boldsymbol{\beta}}] = \boldsymbol{\beta}$ for *every* $\boldsymbol{\beta} \in \mathbb{R}^p$. Since $\boldsymbol{y} = \boldsymbol{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$:

$$\begin{aligned} E[\tilde{\boldsymbol{\beta}}] &= \boldsymbol{C}\, E[\boldsymbol{y}] = \boldsymbol{C}\boldsymbol{X}\boldsymbol{\beta} = [(\boldsymbol{X}^\top \boldsymbol{X})^{-1}\boldsymbol{X}^\top + \boldsymbol{D}]\boldsymbol{X}\boldsymbol{\beta} \\ &= \boldsymbol{I}_p\boldsymbol{\beta} + \boldsymbol{D}\boldsymbol{X}\boldsymbol{\beta} = (\boldsymbol{I}_p + \boldsymbol{D}\boldsymbol{X})\boldsymbol{\beta}. \end{aligned}$$

For this to equal $\boldsymbol{\beta}$ for every $\boldsymbol{\beta}$, we need

$$\boxed{\boldsymbol{D}\boldsymbol{X} = \boldsymbol{0}.} \tag{4}$$

This is the key constraint: the rows of $\boldsymbol{D}$ must be orthogonal to the columns of $\boldsymbol{X}$. In other words, $\boldsymbol{D}$ maps the column space $\mathcal{C}(\boldsymbol{X})$ to zero.

Step 3: Compute the variance and decompose. The variance of $\tilde{\boldsymbol{\beta}}$ is

$$\begin{aligned} \mathrm{Var}(\tilde{\boldsymbol{\beta}}) &= \boldsymbol{C}\,\mathrm{Var}(\boldsymbol{y})\,\boldsymbol{C}^\top = \sigma^2\boldsymbol{C}\boldsymbol{C}^\top \\ &= \sigma^2[(\boldsymbol{X}^\top \boldsymbol{X})^{-1}\boldsymbol{X}^\top + \boldsymbol{D}]\,[\boldsymbol{X}(\boldsymbol{X}^\top \boldsymbol{X})^{-1} + \boldsymbol{D}^\top]. \end{aligned}$$

Expanding the product:

$$\begin{aligned} \boldsymbol{C}\boldsymbol{C}^\top &= (\boldsymbol{X}^\top \boldsymbol{X})^{-1}\underbrace{\boldsymbol{X}^\top \boldsymbol{X}}(\boldsymbol{X}^\top \boldsymbol{X})^{-1} + (\boldsymbol{X}^\top \boldsymbol{X})^{-1}\boldsymbol{X}^\top \boldsymbol{D}^\top + \boldsymbol{D}\boldsymbol{X}(\boldsymbol{X}^\top \boldsymbol{X})^{-1} + \boldsymbol{D}\boldsymbol{D}^\top \\ &= (\boldsymbol{X}^\top \boldsymbol{X})^{-1} + (\boldsymbol{X}^\top \boldsymbol{X})^{-1}\underbrace{(\boldsymbol{D}\boldsymbol{X})^\top}_{\boldsymbol{0}} + \underbrace{\boldsymbol{D}\boldsymbol{X}}_{\boldsymbol{0}}(\boldsymbol{X}^\top \boldsymbol{X})^{-1} + \boldsymbol{D}\boldsymbol{D}^\top, \end{aligned}$$

where the cross terms vanish by the constraint $\boldsymbol{D}\boldsymbol{X} = \boldsymbol{0}$. Therefore:

$$\mathrm{Var}(\tilde{\boldsymbol{\beta}}) = \sigma^2(\boldsymbol{X}^\top \boldsymbol{X})^{-1} + \sigma^2\boldsymbol{D}\boldsymbol{D}^\top = \mathrm{Var}(\hat{\boldsymbol{\beta}}) + \sigma^2\boldsymbol{D}\boldsymbol{D}^\top.$$

Step 4: Conclude. The matrix $\boldsymbol{D}\boldsymbol{D}^\top$ is positive semi-definite (for any vector $\boldsymbol{v}$, $\boldsymbol{v}^\top\boldsymbol{D}\boldsymbol{D}^\top\boldsymbol{v} = \left\|\boldsymbol{D}^\top\boldsymbol{v}\right\|^2 \geq 0$). Therefore

$$\boxed{\mathrm{Var}(\tilde{\boldsymbol{\beta}}) - \mathrm{Var}(\hat{\boldsymbol{\beta}}) = \sigma^2\boldsymbol{D}\boldsymbol{D}^\top \succeq \boldsymbol{0}.}$$

Equality holds if and only if $\boldsymbol{D} = \boldsymbol{0}$, i.e., $\boldsymbol{C} = (\boldsymbol{X}^\top \boldsymbol{X})^{-1}\boldsymbol{X}^\top$. The OLS estimator is the unique Best Linear Unbiased Estimator (BLUE). □

Remark 4.5 (Geometric reading of Gauss–Markov)**.** The constraint $\boldsymbol{D}\boldsymbol{X} = \boldsymbol{0}$ says that the rows of $\boldsymbol{D}$ live in $\mathcal{C}(\boldsymbol{X})^\perp$. Any linear unbiased estimator that differs from OLS must "reach into" the orthogonal complement of the column space to form its estimate. But $\mathcal{C}(\boldsymbol{X})^\perp$ is exactly where the noise lives (the residual space). The extra term $\sigma^2\boldsymbol{D}\boldsymbol{D}^\top$ is the variance penalty for picking up noise from outside the column space. OLS is optimal precisely because it ignores everything orthogonal to $\mathcal{C}(\boldsymbol{X})$.

4.4 R^2 monotonicity

Proposition 4.6. *Adding a predictor to $\boldsymbol{X}$ can only increase (or maintain) R^2.*

Proof. Adding a column to $\boldsymbol{X}$ enlarges $\mathcal{C}(\boldsymbol{X})$. A projection onto a *larger* subspace is at least as close to $\boldsymbol{y}$ as a projection onto a smaller one, so $\|\boldsymbol{e}\|^2$ cannot increase, hence $R^2 = 1 - \|\boldsymbol{e}\|^2/\mathrm{SST}$ cannot decrease. □

Warning

R^2 monotonicity is a geometric fact about the *training sample*. It says nothing about out-of-sample prediction. Adding polynomial terms $(x, x^2, \ldots, x^{12})$ drives R^2 toward 1 while destroying predictive performance—the shadow gets closer in this room but means nothing for the next room.

4.5 Adjusted R^2: paying for dimensions

The monotonicity of R^2 (Proposition 4.6) means that R^2 can never decrease when we add predictors, even useless ones. This makes raw R^2 unreliable for model comparison. The **adjusted** R^2 corrects for this by penalising model complexity.

Definition 4.7 (Adjusted R^2)**.**

$$R^2_{\text{adj}} = 1 - \frac{\text{SSE}/(n-p)}{\text{SST}/(n-1)} = 1 - \frac{n-1}{n-p}(1-R^2).$$

Geometric reading. Think of the data vector $\boldsymbol{y}$ living in $\mathbb{R}^n$. The column space $\mathcal{C}(\boldsymbol{X})$ has dimension p, and the residual space $\mathcal{C}(\boldsymbol{X})^\perp$ has dimension $n-p$. When we project $\boldsymbol{y}$ onto a p-dimensional subspace, we are "using up" p dimensions to capture signal. Even pure noise, projected onto a p-dimensional subspace, produces a nonzero shadow—on average, the squared length of this shadow is $p\sigma^2$.

The residual $\boldsymbol{e}$ lives in the remaining $n-p$ dimensions. If the model captures no real signal, we expect $\|\boldsymbol{e}\|^2 \approx (n-p)\sigma^2$ (noise spread over $n-p$ dimensions) and SST $\approx (n-1)\sigma^2$ (noise spread over $n-1$ dimensions, after demeaning). The ratio $\frac{\text{SSE}/(n-p)}{\text{SST}/(n-1)}$ then hovers around 1, giving $R^2_{\text{adj}} \approx 0$—correctly indicating no explanatory power.

In short, R^2_{adj} accounts for the fact that *projecting onto a subspace costs degrees of freedom*: each dimension consumed by the model must "earn its keep" by reducing the residual more than random chance would.

4.6 Degrees of freedom: the dimension of the residual space

The phrase "degrees of freedom" appears throughout statistics, often without geometric motivation. The projection framework makes it transparent.

Definition 4.8 (Degrees of freedom of the residual)**.** In a linear model with p parameters fitted to n observations, the **residual degrees of freedom** is $n-p$.

Geometric interpretation. The data vector $\boldsymbol{y}$ lives in $\mathbb{R}^n$. The projection $\hat{\boldsymbol{y}} = \boldsymbol{H}\boldsymbol{y}$ lives in $\mathcal{C}(\boldsymbol{X})$, which has dimension p. The residual $\boldsymbol{e} = \boldsymbol{M}\boldsymbol{y}$ lives in $\mathcal{C}(\boldsymbol{X})^\perp$, which has dimension

$$\dim(\mathcal{C}(\boldsymbol{X})^\perp) = n - \dim(\mathcal{C}(\boldsymbol{X})) = n - p.$$

This is confirmed algebraically: $\text{rank}(\boldsymbol{M}) = \text{tr}(\boldsymbol{M}) = \text{tr}(\boldsymbol{I}_n - \boldsymbol{H}) = n - p$.

The residual vector $\boldsymbol{e} \in \mathbb{R}^n$ has n components, but it satisfies p linear constraints ($\boldsymbol{X}^\top\boldsymbol{e} = \boldsymbol{0}$, i.e., $\boldsymbol{e} \perp \mathcal{C}(\boldsymbol{X})$). So $\boldsymbol{e}$ is free to vary in only $n-p$ dimensions. That is the number of "degrees of freedom" left after fitting the model.

Remark 4.9 (Why divide by $n-p$, not n)**.** The unbiased estimator of σ^2 is

$$s^2 = \frac{\|\boldsymbol{e}\|^2}{n-p} = \frac{\text{SSE}}{n-p},$$

not $\|\boldsymbol{e}\|^2/n$. The geometric reason: $\boldsymbol{e} = \boldsymbol{M}\boldsymbol{\varepsilon}$ (since $\boldsymbol{M}\boldsymbol{X} = \boldsymbol{0}$), and the expected squared length of $\boldsymbol{\varepsilon}$ projected onto a $(n-p)$-dimensional subspace is

$$E[\|\boldsymbol{M}\boldsymbol{\varepsilon}\|^2] = E[\boldsymbol{\varepsilon}^\top\boldsymbol{M}\boldsymbol{\varepsilon}] = \sigma^2\,\text{tr}(\boldsymbol{M}) = (n-p)\sigma^2.$$

(Here we used $E[\boldsymbol{\varepsilon}\boldsymbol{\varepsilon}^\top] = \sigma^2\boldsymbol{I}$ and the trace-expectation identity.) Dividing by $n-p$ exactly compensates for the dimensionality of the space into which the noise is projected, yielding $E[s^2] = \sigma^2$.

Dividing by n would systematically underestimate σ^2, because p dimensions of the noise have been "absorbed" by the projection onto $\mathcal{C}(\boldsymbol{X})$—the hat matrix captures p dimensions worth of noise as spurious fit.

So far, we have treated the column space as a monolithic object and projected $\boldsymbol{y}$ onto it in one step. But in multiple regression, different predictors contribute different pieces of information. Can we tease apart the contributions? What does it really mean to "control for" a variable? The Frisch–Waugh–Lovell theorem answers this by decomposing the single projection into sequential steps.

5 Peeling Apart the Variables: Frisch–Waugh–Lovell

5.1 What "controlling for" means geometrically

What does "conditioning on" or "controlling for" a variable mean?

In everyday regression language, you will hear phrases like "the effect of education on income, *controlling for* age" or "the coefficient of x_2 *conditional on* x_1." For an undergraduate, this can sound mysterious. Here is the plain-English idea:

Suppose you want to know whether education affects income. But older people tend to have both more education and higher income simply because they have been working longer. If you just regress income on education, you cannot tell whether the apparent effect comes from education itself or from age lurking in the background. *Controlling for age* means: compare people *of the same age* and then ask whether more education is still associated with higher income. In other words, you strip out the age-related variation from both income and education, and then look at what remains.

In a regression with multiple predictors, this is *exactly* what the coefficient of a predictor measures. When we write $\boldsymbol{y} = \boldsymbol{X}_1\boldsymbol{\beta}_1 + \boldsymbol{X}_2\boldsymbol{\beta}_2 + \boldsymbol{e}$, the coefficient $\hat{\boldsymbol{\beta}}_2$ answers: *what is the effect of $\boldsymbol{X}_2$ after removing ("partialling out") the influence of $\boldsymbol{X}_1$?* The Frisch–Waugh–Lovell theorem below makes this operation precise.

Theorem 5.1 (Frisch–Waugh–Lovell)**.** *Partition the design matrix as $\boldsymbol{X} = [\boldsymbol{X}_1 \;\; \boldsymbol{X}_2]$ with conformable coefficient vector $\boldsymbol{\beta} = (\boldsymbol{\beta}_1^\top, \boldsymbol{\beta}_2^\top)^\top$. Let $\boldsymbol{M}_1 = \boldsymbol{I} - \boldsymbol{X}_1(\boldsymbol{X}_1^\top\boldsymbol{X}_1)^{-1}\boldsymbol{X}_1^\top$ be the annihilator for $\mathcal{C}(\boldsymbol{X}_1)$. Then*

$$\hat{\boldsymbol{\beta}}_2 = (\boldsymbol{X}_2^\top\boldsymbol{M}_1\boldsymbol{X}_2)^{-1}\boldsymbol{X}_2^\top\boldsymbol{M}_1\boldsymbol{y}. \tag{5}$$

That is, $\hat{\boldsymbol{\beta}}_2$ from the full regression equals the OLS coefficient from regressing $\boldsymbol{M}_1\boldsymbol{y}$ on $\boldsymbol{M}_1\boldsymbol{X}_2$—the residualised *response on the* residualised *predictors.*

Proof. The normal equations $\boldsymbol{X}^\top\boldsymbol{X}\,\hat{\boldsymbol{\beta}} = \boldsymbol{X}^\top\boldsymbol{y}$ in block form are

$$\begin{pmatrix} \boldsymbol{X}_1^\top\boldsymbol{X}_1 & \boldsymbol{X}_1^\top\boldsymbol{X}_2 \\ \boldsymbol{X}_2^\top\boldsymbol{X}_1 & \boldsymbol{X}_2^\top\boldsymbol{X}_2 \end{pmatrix} \begin{pmatrix} \hat{\boldsymbol{\beta}}_1 \\ \hat{\boldsymbol{\beta}}_2 \end{pmatrix} = \begin{pmatrix} \boldsymbol{X}_1^\top\boldsymbol{y} \\ \boldsymbol{X}_2^\top\boldsymbol{y} \end{pmatrix}. \tag{6}$$

Step 1: Solve the first block for $\hat{\boldsymbol{\beta}}_1$. The first block equation reads

$$\boldsymbol{X}_1^\top\boldsymbol{X}_1\,\hat{\boldsymbol{\beta}}_1 + \boldsymbol{X}_1^\top\boldsymbol{X}_2\,\hat{\boldsymbol{\beta}}_2 = \boldsymbol{X}_1^\top\boldsymbol{y}.$$

Since $\boldsymbol{X}_1$ has full column rank, $\boldsymbol{X}_1^\top\boldsymbol{X}_1$ is invertible, giving

$$\hat{\boldsymbol{\beta}}_1 = (\boldsymbol{X}_1^\top\boldsymbol{X}_1)^{-1}\boldsymbol{X}_1^\top\boldsymbol{y} - (\boldsymbol{X}_1^\top\boldsymbol{X}_1)^{-1}\boldsymbol{X}_1^\top\boldsymbol{X}_2\,\hat{\boldsymbol{\beta}}_2. \tag{7}$$

Step 2: Substitute into the second block. The second block equation reads

$$\boldsymbol{X}_2^\top\boldsymbol{X}_1\,\hat{\boldsymbol{\beta}}_1 + \boldsymbol{X}_2^\top\boldsymbol{X}_2\,\hat{\boldsymbol{\beta}}_2 = \boldsymbol{X}_2^\top\boldsymbol{y}.$$

Substituting (7):

$$\boldsymbol{X}_2^\top\boldsymbol{X}_1[(\boldsymbol{X}_1^\top\boldsymbol{X}_1)^{-1}\boldsymbol{X}_1^\top\boldsymbol{y} - (\boldsymbol{X}_1^\top\boldsymbol{X}_1)^{-1}\boldsymbol{X}_1^\top\boldsymbol{X}_2\,\hat{\boldsymbol{\beta}}_2] + \boldsymbol{X}_2^\top\boldsymbol{X}_2\,\hat{\boldsymbol{\beta}}_2 = \boldsymbol{X}_2^\top\boldsymbol{y}.$$

Expanding and rearranging:

$$\boldsymbol{X}_2^\top \underbrace{\boldsymbol{X}_1(\boldsymbol{X}_1^\top \boldsymbol{X}_1)^{-1}\boldsymbol{X}_1^\top}_{\boldsymbol{I}-\boldsymbol{M}_1} \boldsymbol{y} - \boldsymbol{X}_2^\top \underbrace{\boldsymbol{X}_1(\boldsymbol{X}_1^\top \boldsymbol{X}_1)^{-1}\boldsymbol{X}_1^\top}_{\boldsymbol{I}-\boldsymbol{M}_1} \boldsymbol{X}_2 \hat{\boldsymbol{\beta}}_2 + \boldsymbol{X}_2^\top \boldsymbol{X}_2 \hat{\boldsymbol{\beta}}_2 = \boldsymbol{X}_2^\top \boldsymbol{y}.$$

Since $\boldsymbol{X}_1(\boldsymbol{X}_1^\top \boldsymbol{X}_1)^{-1}\boldsymbol{X}_1^\top = \boldsymbol{I} - \boldsymbol{M}_1$, we can write:

$$\boldsymbol{X}_2^\top(\boldsymbol{I} - \boldsymbol{M}_1)\boldsymbol{y} + [\boldsymbol{X}_2^\top \boldsymbol{X}_2 - \boldsymbol{X}_2^\top(\boldsymbol{I} - \boldsymbol{M}_1)\boldsymbol{X}_2]\hat{\boldsymbol{\beta}}_2 = \boldsymbol{X}_2^\top \boldsymbol{y}.$$

Step 3: Simplify. Moving $\boldsymbol{X}_2^\top(\boldsymbol{I} - \boldsymbol{M}_1)\boldsymbol{y}$ to the right-hand side:

$$\boldsymbol{X}_2^\top \boldsymbol{M}_1 \boldsymbol{X}_2 \hat{\boldsymbol{\beta}}_2 = \boldsymbol{X}_2^\top \boldsymbol{y} - \boldsymbol{X}_2^\top(\boldsymbol{I} - \boldsymbol{M}_1)\boldsymbol{y} = \boldsymbol{X}_2^\top \boldsymbol{M}_1 \boldsymbol{y}.$$

Here we used $\boldsymbol{X}_2^\top \boldsymbol{X}_2 - \boldsymbol{X}_2^\top(\boldsymbol{I} - \boldsymbol{M}_1)\boldsymbol{X}_2 = \boldsymbol{X}_2^\top \boldsymbol{M}_1 \boldsymbol{X}_2$ and $\boldsymbol{X}_2^\top \boldsymbol{y} - \boldsymbol{X}_2^\top(\boldsymbol{I} - \boldsymbol{M}_1)\boldsymbol{y} = \boldsymbol{X}_2^\top \boldsymbol{M}_1 \boldsymbol{y}$.

Since $\boldsymbol{M}_1$ is symmetric and idempotent ($\boldsymbol{M}_1^2 = \boldsymbol{M}_1$, $\boldsymbol{M}_1^\top = \boldsymbol{M}_1$), we have

$$\boldsymbol{X}_2^\top \boldsymbol{M}_1 \boldsymbol{X}_2 = \boldsymbol{X}_2^\top \boldsymbol{M}_1^\top \boldsymbol{M}_1 \boldsymbol{X}_2 = (\boldsymbol{M}_1 \boldsymbol{X}_2)^\top (\boldsymbol{M}_1 \boldsymbol{X}_2),$$

which is the Gram matrix of the residualised predictors $\boldsymbol{M}_1 \boldsymbol{X}_2$. Similarly, $\boldsymbol{X}_2^\top \boldsymbol{M}_1 \boldsymbol{y} = (\boldsymbol{M}_1 \boldsymbol{X}_2)^\top (\boldsymbol{M}_1 \boldsymbol{y})$.

Therefore, provided $\boldsymbol{M}_1 \boldsymbol{X}_2$ has full column rank (i.e., $\boldsymbol{X}_2$ is not in $\mathcal{C}(\boldsymbol{X}_1)$),

$$\boxed{\hat{\boldsymbol{\beta}}_2 = [(\boldsymbol{M}_1 \boldsymbol{X}_2)^\top (\boldsymbol{M}_1 \boldsymbol{X}_2)]^{-1} (\boldsymbol{M}_1 \boldsymbol{X}_2)^\top (\boldsymbol{M}_1 \boldsymbol{y}),}$$

which is exactly the OLS estimator from regressing the residualised response $\boldsymbol{M}_1 \boldsymbol{y}$ on the residualised predictors $\boldsymbol{M}_1 \boldsymbol{X}_2$. □

Key Geometric Insight

"Controlling for $\boldsymbol{X}_1$" is a geometric operation: project $\boldsymbol{y}$ and $\boldsymbol{X}_2$ onto the orthogonal complement of $\mathcal{C}(\boldsymbol{X}_1)$, then perform simple regression in that cleaned-up space. The multiple regression coefficient is the simple regression coefficient in the residualised subspace.

The geometric analogy: two-step projection. Picture a three-dimensional room. The data vector $\boldsymbol{y}$ is a point floating in the room. The full column space $\mathcal{C}(\boldsymbol{X}_1, \boldsymbol{X}_2)$ is a plane (the "floor"). The subspace $\mathcal{C}(\boldsymbol{X}_1)$ is a line within that floor.

The FWL theorem says you can find the regression coefficient $\hat{\boldsymbol{\beta}}_2$ via two sequential projections instead of one joint projection:

1. **Step 1: Remove $\boldsymbol{X}_1$'s influence.** Project both $\boldsymbol{y}$ and $\boldsymbol{X}_2$ onto the space *perpendicular* to $\mathcal{C}(\boldsymbol{X}_1)$. This is like "looking at the room from directly along the $\boldsymbol{X}_1$ line"—the line collapses to a point, and only the variation perpendicular to $\boldsymbol{X}_1$ remains visible. The residualised vectors $\boldsymbol{M}_1 \boldsymbol{y}$ and $\boldsymbol{M}_1 \boldsymbol{X}_2$ live in this reduced space.
2. **Step 2: Regress in the cleaned space.** Now perform a simple regression of $\boldsymbol{M}_1 \boldsymbol{y}$ on $\boldsymbol{M}_1 \boldsymbol{X}_2$ in the reduced space. The resulting coefficient is exactly $\hat{\boldsymbol{\beta}}_2$ from the full multiple regression.

This two-step procedure is not an approximation—it gives the *exact same* coefficient as the full regression. Geometrically, the full projection onto the plane can be decomposed into a projection onto the line, followed by a projection of the remainder onto the perpendicular direction within the plane.

Example 5.2 (FWL by Hand)**.** We have $n = 4$ observations with two predictors:

i	x_{1i}	x_{2i}	y_i
1	1	0	2
2	0	1	3
3	1	1	6
4	0	0	1

Model: $y_i = \beta_0 + \beta_1 x_{1i} + \beta_2 x_{2i} + e_i$. Notice that x_1 and x_2 have sample means $\bar{x}_1 = \bar{x}_2 = 1/2$ and sample covariance $-1/12$ (moderately negatively correlated among these four points).

Part A: Full regression.

$$\boldsymbol{X} = \begin{pmatrix} 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 1 \\ 1 & 0 & 0 \end{pmatrix}, \quad \boldsymbol{y} = \begin{pmatrix} 2 \\ 3 \\ 6 \\ 1 \end{pmatrix}.$$

$$\boldsymbol{X}^\top \boldsymbol{X} = \begin{pmatrix} 4 & 2 & 2 \\ 2 & 2 & 1 \\ 2 & 1 & 2 \end{pmatrix}, \quad \boldsymbol{X}^\top \boldsymbol{y} = \begin{pmatrix} 12 \\ 8 \\ 9 \end{pmatrix}.$$

Row reduction of the augmented matrix:

$$\left(\begin{array}{ccc|c} 4 & 2 & 2 & 12 \\ 2 & 2 & 1 & 8 \\ 2 & 1 & 2 & 9 \end{array}\right) \xrightarrow{R_2 - \frac{1}{2}R_1,\ R_3 - \frac{1}{2}R_1} \left(\begin{array}{ccc|c} 4 & 2 & 2 & 12 \\ 0 & 1 & 0 & 2 \\ 0 & 0 & 1 & 3 \end{array}\right).$$

Reading off: $\hat{\beta}_2 = 3$, $\hat{\beta}_1 = 2$, $\hat{\beta}_0 = (12 - 2 \cdot 2 - 2 \cdot 3)/4 = (12 - 4 - 6)/4 = 1/2$.

So $\hat{\boldsymbol{\beta}} = (1/2,\ 2,\ 3)^\top$, and in particular $\boxed{\hat{\beta}_2 = 3}$.

Verification:

$$\hat{\boldsymbol{y}} = \begin{pmatrix} 1/2 + 2 + 0 \\ 1/2 + 0 + 3 \\ 1/2 + 2 + 3 \\ 1/2 + 0 + 0 \end{pmatrix} = \begin{pmatrix} 5/2 \\ 7/2 \\ 11/2 \\ 1/2 \end{pmatrix}, \quad \boldsymbol{e} = \begin{pmatrix} 2 - 5/2 \\ 3 - 7/2 \\ 6 - 11/2 \\ 1 - 1/2 \end{pmatrix} = \begin{pmatrix} -1/2 \\ -1/2 \\ 1/2 \\ 1/2 \end{pmatrix}.$$

Check: $\boldsymbol{X}^\top \boldsymbol{e} = (-1/2 - 1/2 + 1/2 + 1/2,\ -1/2 + 1/2,\ -1/2 + 1/2)^\top = (0, 0, 0)^\top$. ✓

Part B: FWL two-step procedure for $\hat{\beta}_2$.

Partition $\boldsymbol{X}_1 = (\mathbf{1} \mid \boldsymbol{x}_1)$ (intercept and first predictor), $\boldsymbol{X}_2 = \boldsymbol{x}_2 = (0, 1, 1, 0)^\top$.

Step B1: Compute $(\boldsymbol{X}_1^\top \boldsymbol{X}_1)^{-1}$.

$$\boldsymbol{X}_1 = \begin{pmatrix} 1 & 1 \\ 1 & 0 \\ 1 & 1 \\ 1 & 0 \end{pmatrix}, \quad \boldsymbol{X}_1^\top \boldsymbol{X}_1 = \begin{pmatrix} 4 & 2 \\ 2 & 2 \end{pmatrix}, \quad \det = 4, \quad (\boldsymbol{X}_1^\top \boldsymbol{X}_1)^{-1} = \frac{1}{4}\begin{pmatrix} 2 & -2 \\ -2 & 4 \end{pmatrix} = \begin{pmatrix} 1/2 & -1/2 \\ -1/2 & 1 \end{pmatrix}.$$

Step B2: Residualise $\boldsymbol{x}_2$ *on* $\boldsymbol{X}_1$.

Regress $\boldsymbol{x}_2 = (0, 1, 1, 0)^\top$ on $\boldsymbol{X}_1$:

$$\boldsymbol{X}_1^\top \boldsymbol{x}_2 = \begin{pmatrix} 2 \\ 1 \end{pmatrix}, \quad (\boldsymbol{X}_1^\top \boldsymbol{X}_1)^{-1} \boldsymbol{X}_1^\top \boldsymbol{x}_2 = \begin{pmatrix} 1/2 & -1/2 \\ -1/2 & 1 \end{pmatrix} \begin{pmatrix} 2 \\ 1 \end{pmatrix} = \begin{pmatrix} 1/2 \\ 0 \end{pmatrix}.$$

Projection: $\boldsymbol{X}_1 \left(\begin{smallmatrix} 1/2 \\ 0 \end{smallmatrix}\right) = (1/2,\ 1/2,\ 1/2,\ 1/2)^\top$.

Residual:

$$\tilde{\boldsymbol{x}}_2 = (0, 1, 1, 0)^\top - (1/2, 1/2, 1/2, 1/2)^\top = (-1/2,\ 1/2,\ 1/2,\ -1/2)^\top.$$

This is the part of x_2 not explained by the intercept and x_1—the "new information" that x_2 brings beyond what x_1 already provides.

Step B3: Residualise $\boldsymbol{y}$ *on* $\boldsymbol{X}_1$.

$$\boldsymbol{X}_1^\top \boldsymbol{y} = \begin{pmatrix} 12 \\ 8 \end{pmatrix}, \quad (\boldsymbol{X}_1^\top \boldsymbol{X}_1)^{-1} \boldsymbol{X}_1^\top \boldsymbol{y} = \begin{pmatrix} 1/2 & -1/2 \\ -1/2 & 1 \end{pmatrix} \begin{pmatrix} 12 \\ 8 \end{pmatrix} = \begin{pmatrix} 2 \\ 2 \end{pmatrix}.$$

Projection: $\boldsymbol{X}_1\binom{2}{2} = (2+2,\ 2+0,\ 2+2,\ 2+0)^\top = (4,\ 2,\ 4,\ 2)^\top$.

Residual:

$$\tilde{\boldsymbol{y}} = (2,3,6,1)^\top - (4,2,4,2)^\top = (-2,\ 1,\ 2,\ -1)^\top.$$

Step B4: Regress $\tilde{\boldsymbol{y}}$ on $\tilde{\boldsymbol{x}}_2$ (no intercept needed—both have been purged of $\boldsymbol{X}_1$).

$$\hat{\beta}_2^{\text{FWL}} = \frac{\tilde{\boldsymbol{x}}_2^\top \tilde{\boldsymbol{y}}}{\tilde{\boldsymbol{x}}_2^\top \tilde{\boldsymbol{x}}_2}.$$

Numerator:

$$\tilde{\boldsymbol{x}}_2^\top \tilde{\boldsymbol{y}} = (-\tfrac{1}{2})(-2) + (\tfrac{1}{2})(1) + (\tfrac{1}{2})(2) + (-\tfrac{1}{2})(-1) = 1 + \tfrac{1}{2} + 1 + \tfrac{1}{2} = 3.$$

Denominator:

$$\tilde{\boldsymbol{x}}_2^\top \tilde{\boldsymbol{x}}_2 = \tfrac{1}{4} + \tfrac{1}{4} + \tfrac{1}{4} + \tfrac{1}{4} = 1.$$

Therefore:

$$\boxed{\hat{\beta}_2^{\text{FWL}} = \frac{3}{1} = 3 = \hat{\beta}_2^{\text{full}}.} \checkmark$$

What happened geometrically? The FWL theorem says: to find the coefficient of x_2 in the full model, first remove from both $\boldsymbol{y}$ and $\boldsymbol{x}_2$ everything that $\boldsymbol{X}_1$ can explain. Regressing the residual components on each other gives *exactly* the same coefficient as the full regression.

Takeaway. FWL is not an approximation. The two-step residualisation procedure recovers $\hat{\beta}_2 = 3$ exactly, because projection onto a direct sum decomposes into sequential projections onto complementary subspaces.

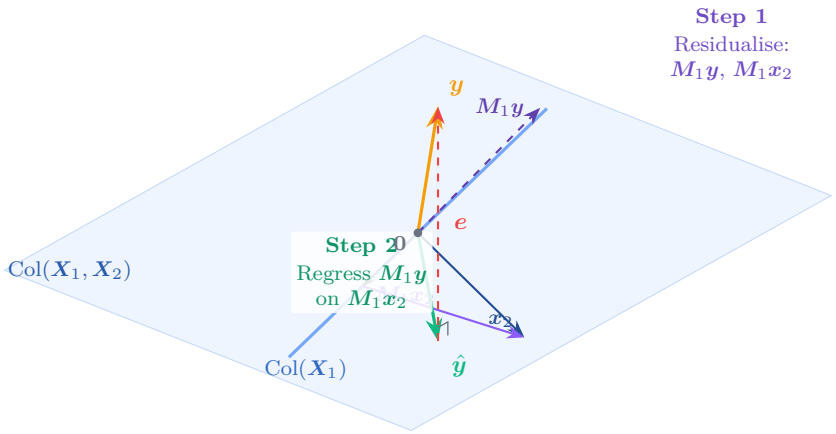

Figure 4: The Frisch–Waugh–Lovell theorem as a two-step projection. **Step 1**: residualise both $\boldsymbol{y}$ and $\boldsymbol{x}_2$ with respect to $\mathrm{Col}(\boldsymbol{X}_1)$ using the annihilator $\boldsymbol{M}_1$, sweeping out the effect of $\boldsymbol{X}_1$. **Step 2**: regress $\boldsymbol{M}_1\boldsymbol{y}$ on $\boldsymbol{M}_1\boldsymbol{x}_2$ in the reduced orthogonal complement. The resulting coefficient $\hat{\beta}_2$ is identical to the one obtained from the full regression, and the final fitted vector $\hat{\boldsymbol{y}}$ is the same.

? Stop and Think

Two limiting cases of FWL.
Recall that $\hat{\boldsymbol{\beta}}_2 = (\boldsymbol{X}_2^\top \boldsymbol{M}_1 \boldsymbol{X}_2)^{-1} \boldsymbol{X}_2^\top \boldsymbol{M}_1 \boldsymbol{y}$, where $\boldsymbol{M}_1 = \boldsymbol{I} - \boldsymbol{H}_1$ removes the influence of $\boldsymbol{X}_1$.

1. **Case 1:** $\boldsymbol{X}_1 \perp \boldsymbol{X}_2$ (i.e. $\boldsymbol{X}_1^\top \boldsymbol{X}_2 = \boldsymbol{0}$). What does $\boldsymbol{M}_1 \boldsymbol{X}_2$ simplify to? What does this tell you about the FWL coefficient $\hat{\boldsymbol{\beta}}_2$ compared to the coefficient from a simple regression of $\boldsymbol{y}$ on $\boldsymbol{X}_2$ alone?
2. **Case 2:** $\boldsymbol{X}_2 \in \mathcal{C}(\boldsymbol{X}_1)$ (i.e. $\boldsymbol{X}_2$ is a linear combination of the columns of $\boldsymbol{X}_1$). What does $\boldsymbol{M}_1 \boldsymbol{X}_2$ simplify to? What goes wrong?

Answers: (1) When $\boldsymbol{X}_1 \perp \boldsymbol{X}_2$, we have $\boldsymbol{H}_1 \boldsymbol{X}_2 = \boldsymbol{0}$, so $\boldsymbol{M}_1 \boldsymbol{X}_2 = \boldsymbol{X}_2$. The residualised predictor equals the original predictor—controlling for $\boldsymbol{X}_1$ changes nothing. The FWL coefficient equals the simple regression coefficient. This is why uncorrelated predictors give the same coefficients whether you include them together or separately.
(2) When $\boldsymbol{X}_2 \in \mathcal{C}(\boldsymbol{X}_1)$, we have $\boldsymbol{H}_1 \boldsymbol{X}_2 = \boldsymbol{X}_2$, so $\boldsymbol{M}_1 \boldsymbol{X}_2 = \boldsymbol{0}$. The residualised predictor is the zero vector. The formula $(\boldsymbol{X}_2^\top \boldsymbol{M}_1 \boldsymbol{X}_2)^{-1}$ requires inverting a zero matrix, which is impossible. This is perfect multicollinearity: $\boldsymbol{X}_2$ contains no information beyond what $\boldsymbol{X}_1$ already provides.

5.2 The effect of predictor correlation

The FWL theorem reveals *why* correlated predictors cause trouble. To understand this, think carefully about what the residualisation step does.

The residualised predictor and its length. When we "control for" $\boldsymbol{X}_1$, we replace $\boldsymbol{X}_2$ by its residuals after regressing on $\boldsymbol{X}_1$:

$$\boldsymbol{M}_1 \boldsymbol{X}_2 = \boldsymbol{X}_2 - \underbrace{\boldsymbol{X}_1 (\boldsymbol{X}_1^\top \boldsymbol{X}_1)^{-1} \boldsymbol{X}_1^\top \boldsymbol{X}_2}_{\text{part of } \boldsymbol{X}_2 \text{ explained by } \boldsymbol{X}_1}.$$

The residualised predictor $\boldsymbol{M}_1 \boldsymbol{X}_2$ contains only the "unique information" in $\boldsymbol{X}_2$—the part that $\boldsymbol{X}_1$ cannot account for.

Now think about what happens as the correlation between $\boldsymbol{X}_1$ and $\boldsymbol{X}_2$ changes:

- **Uncorrelated predictors** ($\boldsymbol{X}_1^\top \boldsymbol{X}_2 \approx \boldsymbol{0}$): The projection of $\boldsymbol{X}_2$ onto $\mathcal{C}(\boldsymbol{X}_1)$ is nearly zero, so $\boldsymbol{M}_1 \boldsymbol{X}_2 \approx \boldsymbol{X}_2$. Almost all of $\boldsymbol{X}_2$'s information is unique. The residualised predictor is long, and the coefficient is estimated precisely.
- **Moderately correlated predictors**: Part of $\boldsymbol{X}_2$ is "used up" by $\boldsymbol{X}_1$. The residualised predictor $\boldsymbol{M}_1 \boldsymbol{X}_2$ is shorter. Less unique information remains, and the coefficient estimate becomes noisier.
- **Highly correlated predictors** ($\boldsymbol{X}_2 \approx \boldsymbol{X}_1 \boldsymbol{c}$ for some $\boldsymbol{c}$): Nearly all of $\boldsymbol{X}_2$ is explained by $\boldsymbol{X}_1$. The residualised predictor $\boldsymbol{M}_1 \boldsymbol{X}_2$ is very short—close to the zero vector. There is almost no unique information left. The coefficient estimate becomes wildly unstable.

Geometric picture. Imagine $\boldsymbol{X}_1$ and $\boldsymbol{X}_2$ as two vectors in $\mathbb{R}^n$. Residualising $\boldsymbol{X}_2$ with respect to $\boldsymbol{X}_1$ means projecting $\boldsymbol{X}_2$ onto the plane perpendicular to $\boldsymbol{X}_1$ and keeping only the perpendicular component. If $\boldsymbol{X}_2$ points in nearly the same direction as $\boldsymbol{X}_1$, this perpendicular component is tiny. You are trying to estimate a coefficient by regressing on a near-zero vector, which amplifies any noise enormously.

The standard error formula. The standard error of $\hat{\boldsymbol{\beta}}_2$ is

$$\mathrm{SE}(\hat{\beta}_{2,j}) = \sqrt{\frac{s^2}{\|\boldsymbol{M}_1 \boldsymbol{x}_{2,j}\|^2}},$$

where $s^2 = \|\boldsymbol{e}\|^2 / (n-p)$ is the estimated residual variance. As correlation increases, $\|\boldsymbol{M}_1 \boldsymbol{x}_{2,j}\| \to 0$ and the denominator vanishes, sending SE $\to \infty$. This is the geometric root of *multicollinearity*: the coefficient is not wrong on average (it is still unbiased), but it is estimated so imprecisely that it becomes practically useless.

FWL revealed that individual coefficients depend on the unique information each predictor contributes, as measured by the residualised predictor's length. But some observations contribute more to the projection than others. Which observations have the most influence on the fitted values, and how can we detect problematic data points?